

```
# import all libraries and dependencies for dataframe
```

```
import pandas as pd
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
#loading data
```

```
df = pd.read_csv("D:\Data Visualization(IShant)\
```

```
CarPrice_Assignment.csv")
```

```
df
```

	car_ID	symboling	CarName	fueltype	
aspiration \					
0	1	3	alfa-romero giulia	gas	std
1	2	3	alfa-romero stelvio	gas	std
2	3	1	alfa-romero Quadrifoglio	gas	std
3	4	2	audi 100 ls	gas	std
4	5	2	audi 100ls	gas	std
..	...	...	...	...	...
200	201	-1	volvo 145e (sw)	gas	std
201	202	-1	volvo 144ea	gas	turbo
202	203	-1	volvo 244dl	gas	std
203	204	-1	volvo 246	diesel	turbo
204	205	-1	volvo 264gl	gas	turbo

	doornumber	carbody	drivewheel	engine	location	wheelbase	...
\							
0	two	convertible	rwd	front	88.6	...	
1	two	convertible	rwd	front	88.6	...	
2	two	hatchback	rwd	front	94.5	...	
3	four	sedan	fwd	front	99.8	...	
4	four	sedan	4wd	front	99.4	...	
..	...	...	...	...	...	...	
200	four	sedan	rwd	front	109.1	...	

201	four	sedan	rwd	front	109.1	...
202	four	sedan	rwd	front	109.1	...
203	four	sedan	rwd	front	109.1	...
204	four	sedan	rwd	front	109.1	...

	engine size	fuel system	bore ratio	stroke	compression ratio
0	130	mpfi	3.47	2.68	9.0
1	130	mpfi	3.47	2.68	9.0
2	152	mpfi	2.68	3.47	9.0
3	109	mpfi	3.19	3.40	10.0
4	136	mpfi	3.19	3.40	8.0
..	...	...	...	...	...
200	141	mpfi	3.78	3.15	9.5
201	141	mpfi	3.78	3.15	8.7
202	173	mpfi	3.58	2.87	8.8
203	145	idi	3.01	3.40	23.0
204	141	mpfi	3.78	3.15	9.5

	peak rpm	city mpg	highway mpg	price
0	5000	21	27	13495.0
1	5000	21	27	16500.0
2	5000	19	26	16500.0
3	5500	24	30	13950.0
4	5500	18	22	17450.0
..	...	...	...	...
200	5400	23	28	16845.0
201	5300	19	25	19045.0
202	5500	18	23	21485.0
203	4800	26	27	22470.0
204	5400	19	25	22625.0

[205 rows x 26 columns]

```
## information of the data
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 205 entries, 0 to 204
```

```
Data columns (total 26 columns):
```

#	Column	Non-Null Count	Dtype
0	car_ID	205 non-null	int64
1	symboling	205 non-null	int64
2	CarName	205 non-null	object
3	fueltype	205 non-null	object
4	aspiration	205 non-null	object
5	doornumber	205 non-null	object
6	carbody	205 non-null	object
7	drivewheel	205 non-null	object
8	enginelocation	205 non-null	object
9	wheelbase	205 non-null	float64
10	carlength	205 non-null	float64
11	carwidth	205 non-null	float64
12	carheight	205 non-null	float64
13	curbweight	205 non-null	int64
14	enginetype	205 non-null	object
15	cylindernumber	205 non-null	object
16	enginesize	205 non-null	int64
17	fuelsystem	205 non-null	object
18	boreratio	205 non-null	float64
19	stroke	205 non-null	float64
20	compressionratio	205 non-null	float64
21	horsepower	205 non-null	int64
22	peakrpm	205 non-null	int64
23	citympg	205 non-null	int64
24	highwaympg	205 non-null	int64
25	price	205 non-null	float64

```
dtypes: float64(8), int64(8), object(10)
```

```
memory usage: 41.8+ KB
```

```
# unique value of each column
```

```
df.nunique()
```

car_ID	205
symboling	6
CarName	147
fueltype	2
aspiration	2
doornumber	2
carbody	5
drivewheel	3
enginelocation	2
wheelbase	53

```

carlength      75
carwidth       44
carheight      49
curbweight     171
enginetype      7
cylindernumber  7
enginesize     44
fuelsystem      8
boreratio      38
stroke         37
compressionratio 32
horsepower     59
peakrpm        23
citympg        29
highwaympg     30
price         189
dtype: int64

```

```

# description of the data
df.describe()

```

```

      car_ID  symboling  wheelbase  carlength  carwidth
carheight \
count  205.000000  205.000000  205.000000  205.000000  205.000000
205.000000
mean   103.000000   0.834146   98.756585  174.049268   65.907805
53.724878
std     59.322565   1.245307   6.021776   12.337289   2.145204
2.443522
min      1.000000  -2.000000   86.600000  141.100000   60.300000
47.800000
25%     52.000000   0.000000   94.500000  166.300000   64.100000
52.000000
50%     103.000000  1.000000   97.000000  173.200000   65.500000
54.100000
75%     154.000000  2.000000  102.400000  183.100000   66.900000
55.500000
max     205.000000  3.000000  120.900000  208.100000   72.300000
59.800000

      curbweight  enginesize  boreratio  stroke
compressionratio \
count  205.000000  205.000000  205.000000  205.000000
205.000000
mean   2555.565854  126.907317   3.329756   3.255415
10.142537
std     520.680204  41.642693   0.270844   0.313597
3.972040
min   1488.000000  61.000000   2.540000   2.070000
7.000000

```

25%	2145.000000	97.000000	3.150000	3.110000
8.600000				
50%	2414.000000	120.000000	3.310000	3.290000
9.000000				
75%	2935.000000	141.000000	3.580000	3.410000
9.400000				
max	4066.000000	326.000000	3.940000	4.170000
23.000000				

	horsepower	peakrpm	citympg	highwaympg	price
count	205.000000	205.000000	205.000000	205.000000	205.000000
mean	104.117073	5125.121951	25.219512	30.751220	13276.710571
std	39.544167	476.985643	6.542142	6.886443	7988.852332
min	48.000000	4150.000000	13.000000	16.000000	5118.000000
25%	70.000000	4800.000000	19.000000	25.000000	7788.000000
50%	95.000000	5200.000000	24.000000	30.000000	10295.000000
75%	116.000000	5500.000000	30.000000	34.000000	16503.000000
max	288.000000	6600.000000	49.000000	54.000000	45400.000000

```
# Dropping car_ID from the database
```

```
df=df.drop(['car_ID'],axis=1)
```

```
# Extracting Car Company from the CarName
```

```
df['CarName']=df['CarName'].str.split(' ').str[0]
```

```
# Unique Car company
```

```
df['CarName'].unique()
```

```
array(['alfa-romero', 'audi', 'bmw', 'chevrolet', 'dodge', 'honda',
      'isuzu', 'jaguar', 'maxda', 'mazda', 'buick', 'mercury',
      'mitsubishi', 'Nissan', 'nissan', 'peugeot', 'plymouth',
      'porsche',
      'porcshce', 'renault', 'saab', 'subaru', 'toyota', 'toyouta',
      'vokswagen', 'volkswagen', 'vw', 'volvo'], dtype=object)
```

Here, type error for car namming have been found for different brand So;

```
# Renaming the typo errors in Car Company names
```

```
df['CarName'] = df['CarName'].replace({'maxda': 'mazda', 'nissan':
'nissan', 'porcshce': 'porsche', 'toyouta': 'toyota',
'vokswagen': 'volkswagen', 'vw':
'volkswagen'})
```

```
# Unique Car company
```

```
df['CarName'].unique()
```

```
array(['alfa-romero', 'audi', 'bmw', 'chevrolet', 'dodge', 'honda',
      'isuzu', 'jaguar', 'mazda', 'buick', 'mercury', 'mitsubishi',
      'Nissan', 'nissan', 'peugeot', 'plymouth', 'porsche',
```

```

'renault',
    'saab', 'subaru', 'toyota', 'volkswagen', 'volvo'],
dtype=object)

# Check for duplicated rows
duplicated_rows_count = df.duplicated().sum()

# Display the count of duplicated rows
print("Duplicated Rows Count:", duplicated_rows_count)

Duplicated Rows Count: 0

# Changing the datatype of symboling as it is categorical variable
df['symboling']=df['symboling'].astype('str')

# Numerical and Categorical Columns
numerical_columns=df.select_dtypes(exclude=['object']).columns
categorical_columns=df.select_dtypes(include=['object']).columns

# Display the lists of numerical and categorical columns
print("Numerical Columns:", numerical_columns)
print("Categorical Columns:", categorical_columns)

Numerical Columns: Index(['wheelbase', 'carlength', 'carwidth',
    'carheight', 'curbweight',
    'enginesize', 'bore ratio', 'stroke', 'compressionratio',
    'horsepower',
    'peakrpm', 'citympg', 'highwaympg', 'price'],
    dtype='object')
Categorical Columns: Index(['symboling', 'CarName', 'fueltype',
    'aspiration', 'doornumber',
    'carbody', 'drivewheel', 'engine location', 'enginetype',
    'cylindernumber', 'fuelsystem'],
    dtype='object')

```

Showing the data nature of index of columns in file

```
df.shape
```

```
(205, 25)
```

```
df.describe().T
```

	count	mean	std	min	max
25% \ wheelbase	205.0	98.756585	6.021776	86.60	94.50
carlength	205.0	174.049268	12.337289	141.10	166.30
carwidth	205.0	65.907805	2.145204	60.30	64.10

carheight	205.0	53.724878	2.443522	47.80	52.00
curbweight	205.0	2555.565854	520.680204	1488.00	2145.00
enginesize	205.0	126.907317	41.642693	61.00	97.00
boreratio	205.0	3.329756	0.270844	2.54	3.15
stroke	205.0	3.255415	0.313597	2.07	3.11
compressionratio	205.0	10.142537	3.972040	7.00	8.60
horsepower	205.0	104.117073	39.544167	48.00	70.00
peakrpm	205.0	5125.121951	476.985643	4150.00	4800.00
citympg	205.0	25.219512	6.542142	13.00	19.00
highwaympg	205.0	30.751220	6.886443	16.00	25.00
price	205.0	13276.710571	7988.852332	5118.00	7788.00

	50%	75%	max
wheelbase	97.00	102.40	120.90
carlength	173.20	183.10	208.10
carwidth	65.50	66.90	72.30
carheight	54.10	55.50	59.80
curbweight	2414.00	2935.00	4066.00
enginesize	120.00	141.00	326.00
boreratio	3.31	3.58	3.94
stroke	3.29	3.41	4.17
compressionratio	9.00	9.40	23.00
horsepower	95.00	116.00	288.00
peakrpm	5200.00	5500.00	6600.00
citympg	24.00	30.00	49.00
highwaympg	30.00	34.00	54.00
price	10295.00	16503.00	45400.00

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 205 entries, 0 to 204
```

```
Data columns (total 25 columns):
```

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	symboling	205 non-null	object
1	CarName	205 non-null	object
2	fueltype	205 non-null	object
3	aspiration	205 non-null	object
4	doornumber	205 non-null	object

```

5   carbody      205 non-null   object
6   drivewheel  205 non-null   object
7   enginelocation 205 non-null   object
8   wheelbase   205 non-null   float64
9   carlength   205 non-null   float64
10  carwidth     205 non-null   float64
11  carheight    205 non-null   float64
12  curbweight   205 non-null   int64
13  enginetype   205 non-null   object
14  cylindernumber 205 non-null   object
15  enginesize    205 non-null   int64
16  fuelsystem   205 non-null   object
17  boreratio    205 non-null   float64
18  stroke       205 non-null   float64
19  compressionratio 205 non-null   float64
20  horsepower    205 non-null   int64
21  peakrpm       205 non-null   int64
22  citympg       205 non-null   int64
23  highwaympg    205 non-null   int64
24  price         205 non-null   float64
dtypes: float64(8), int64(6), object(11)
memory usage: 40.2+ KB

```

```
# Check for missing values in each column
```

```
missing_values_count = df.isnull().sum()
```

```
# Display the count of missing values
```

```
print("Missing Values Count:")
```

```
print(missing_values_count)
```

```

Missing Values Count:
symboling      0
CarName        0
fueltype       0
aspiration     0
doornumber     0
carbody        0
drivewheel     0
enginelocation 0
wheelbase      0
carlength      0
carwidth       0
carheight      0
curbweight     0
enginetype     0
cylindernumber 0
enginesize     0
fuelsystem     0
boreratio      0
stroke         0

```



```

compressionratio    0
horsepower          0
peakrpm             0
citympg             0
highwaympg          0
price               0
dtype: int64

```

It appears that there are no missing values (NaN) in any of the columns of your DataFrame. Each column has a count of 0 missing values.

```

# Get numerical columns
nums_cols = df.select_dtypes(include=["float64", "int64"]).columns

# Calculate the number of rows and columns needed for the subplots
num_plots = len(nums_cols)
num_rows = (num_plots - 1) // 4 + 1
num_cols = min(num_plots, 4)

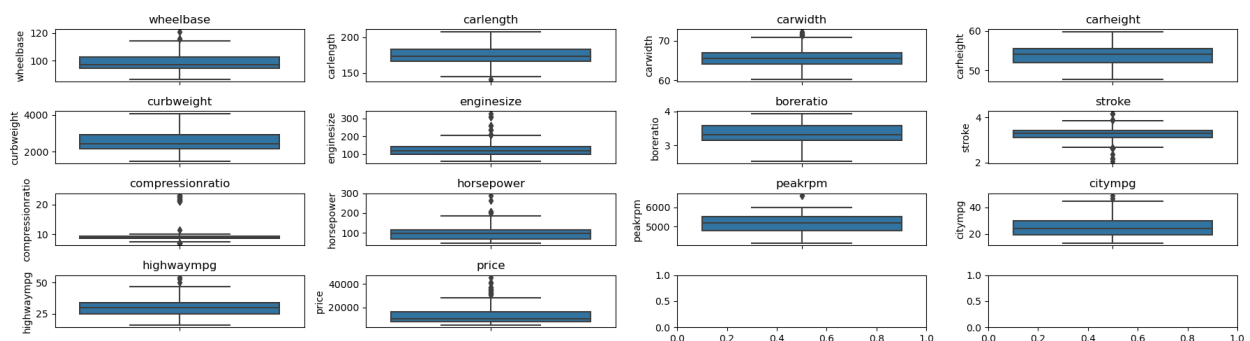
# Create subplots
fig, axes = plt.subplots(num_rows, num_cols, figsize=(18, 5))

# Flatten axes if there's only one row or column
if num_rows == 1 or num_cols == 1:
    axes = axes.reshape(-1)

# Plot boxplots
for i, col in enumerate(nums_cols):
    if num_rows > 1 and num_cols > 1:
        row = i // num_cols
        col_index = i % num_cols
        sns.boxplot(y=df[col], ax=axes[row, col_index])
        axes[row, col_index].set_title(col)
    else:
        sns.boxplot(y=df[col], ax=axes[i])
        axes[i].set_title(col)

plt.tight_layout()
plt.show()


```



Columns Having Outliers :

Wheelbase,Carlength,Carwidth,enginesize,stroke,compressionratio,horsepower,peakrpm,citympg,highwaympg,price

```
# Create a heatmap of the correlation matrix
plt.figure(figsize=(16, 7))
sns.heatmap(df.corr(numeric_only=True), annot=True, cmap='Blues')
plt.title('Correlation Matrix Heatmap')
plt.show()
```



```
# Drop specified columns
columns_to_drop = ['symboling', 'stroke', 'compressionratio',
'peakrpm', 'citympg', 'highwaympg', 'carlength', 'curbweight']
df = df.drop(columns=columns_to_drop, axis=1)
```

Dropping down the these cells because of having high correlation between each other as they stand by independent variable. So, mustnot have to high correlated with each other as explanatory variables

```
df
      CarName fueltype aspiration doornumber   carbody
drivewheel \
0   alfa-romero    gas         std         two  convertible
rwd
1   alfa-romero    gas         std         two  convertible
rwd
2   alfa-romero    gas         std         two   hatchback
```

rwd					
3	audi	gas	std	four	sedan
fwd					
4	audi	gas	std	four	sedan
4wd					
..	...	...	...	...	...
.					
200	volvo	gas	std	four	sedan
rwd					
201	volvo	gas	turbo	four	sedan
rwd					
202	volvo	gas	std	four	sedan
rwd					
203	volvo	diesel	turbo	four	sedan
rwd					
204	volvo	gas	turbo	four	sedan
rwd					
	engine	location	wheelbase	carwidth	carheight
cylinders	number	\			engine
0	front		88.6	64.1	48.8
four					dohc
1	front		88.6	64.1	48.8
four					dohc
2	front		94.5	65.5	52.4
six					ohcv
3	front		99.8	66.2	54.3
four					ohc
4	front		99.4	66.4	54.3
five					ohc
..	...		...	...	...
...					
200	front		109.1	68.9	55.5
four					ohc
201	front		109.1	68.8	55.5
four					ohc
202	front		109.1	68.9	55.5
six					ohcv
203	front		109.1	68.9	55.5
six					ohc
204	front		109.1	68.9	55.5
four					ohc
	engine	size	fuel	system	boreratio
0	130	mpfi			3.47
1	130	mpfi			3.47
2	152	mpfi			2.68
3	109	mpfi			3.19
4	136	mpfi			3.19
				horsepower	price
				111	13495.0
				111	16500.0
				154	16500.0
				102	13950.0
				115	17450.0

```

200      141      mpfi      3.78      114  16845.0
201      141      mpfi      3.78      160  19045.0
202      173      mpfi      3.58      134  21485.0
203      145       idi      3.01      106  22470.0
204      141      mpfi      3.78      114  22625.0


```

```
[205 rows x 17 columns]
```

```

# Create a heatmap of the correlation matrix
plt.figure(figsize=(16, 7))
sns.heatmap(df.corr(numeric_only=True), annot=True, cmap='Blues')
plt.title('Correlation Matrix Heatmap')
plt.show()

```



```
df.head()
```

	CarName	fueltype	aspiration	doornumber	carbody	drivewheel
0	alfa-romero	gas	std	two	convertible	rwd
1	alfa-romero	gas	std	two	convertible	rwd
2	alfa-romero	gas	std	two	hatchback	rwd
3	audi	gas	std	four	sedan	fwd
4	audi	gas	std	four	sedan	4wd

engine	location	wheelbase	carwidth	carheight	enginetype
cylindernumber \					
0	front	88.6	64.1	48.8	dohc
four					
1	front	88.6	64.1	48.8	dohc
four					
2	front	94.5	65.5	52.4	ohcv
six					
3	front	99.8	66.2	54.3	ohc
four					
4	front	99.4	66.4	54.3	ohc
five					

	enginesize	fuelsystem	boreratio	horsepower	price
0	130	mpfi	3.47	111	13495.0
1	130	mpfi	3.47	111	16500.0
2	152	mpfi	2.68	154	16500.0
3	109	mpfi	3.19	102	13950.0
4	136	mpfi	3.19	115	17450.0

```
df['fueltype'].value_counts()
```

```
fueltype
gas      185
diesel   20
Name: count, dtype: int64
```

The output you provided shows that there are 185 occurrences of 'gas' and 20 occurrences of 'diesel'

```
df['fueltype'] = df['fueltype'].replace({'gas':1,'diesel':2})
#transform categorical data into numerical data in the 'fueltype' column
```

```
df['aspiration'].value_counts()
```

```
aspiration
std      168
turbo    37
Name: count, dtype: int64
```

```
df['aspiration'] = df['aspiration'].replace({'std':1,'turbo':2})
```

```
df['doornumber'].value_counts()
```

```
doornumber
four     115
two       90
Name: count, dtype: int64
```

```
df['doornumber'] = df['doornumber'].replace({'four':4,'two':2})
```

```

df['carbody'].value_counts()

carbody
sedan      96
hatchback  70
wagon      25
hardtop     8
convertible 6
Name: count, dtype: int64

df['carbody'] =
df['carbody'].replace(['wagon', 'hardtop', 'convertible'], 'others')
df['carbody'] =
df['carbody'].replace({'sedan':1, 'hatchback':2, 'others':3})

df['drivewheel'].value_counts()

drivewheel
fwd      120
rwd       76
4wd        9
Name: count, dtype: int64

df['drivewheel'] = df['drivewheel'].replace({'fwd':1, 'rwd':2, '4wd':3})

df['enginelocation'].value_counts()

enginelocation
front     202
rear        3
Name: count, dtype: int64

df['enginelocation'] =
df['enginelocation'].replace({'front':1, 'rear':2})

df['enginetype'].value_counts()

enginetype
ohc      148
ohcf      15
ohcv      13
dohc      12
l         12
rotor       4
dohcv       1
Name: count, dtype: int64

df['enginetype'] =
df['enginetype'].replace(['ohcf', 'ohcv', 'dohc', 'l', 'rotor', 'dohcv'], 'others')
df['enginetype'] = df['enginetype'].replace({'ohc':1, 'others':2})

```

```

df['cylindernumber'].value_counts()

cylindernumber
four      159
six       24
five      11
eight      5
two        4
three      1
twelve     1
Name: count, dtype: int64

df['cylindernumber'] =
df['cylindernumber'].replace(['five','eight','two','three','twelve'],'
others')
df['cylindernumber'] =
df['cylindernumber'].replace({'four':4,'six':6,'others':1})

df['fuelsystem'].value_counts()

fuelsystem
mpfi      94
2bbl      66
idi       20
1bbl      11
spdi       9
4bbl       3
mfi        1
spfi       1
Name: count, dtype: int64

df['fuelsystem'] =
df['fuelsystem'].replace(['idi','1bbl','spdi','4bbl','mfi','spfi'],'ot
hers')
df['fuelsystem'] =
df['fuelsystem'].replace({'mpfi':1,'2bbl':2,'others':3})

df.head()

```


	CarName	fueltype	aspiration	doornumber	carbody	drivewheel
0	alfa-romero	1	1	2	3	2
1	alfa-romero	1	1	2	3	2
2	alfa-romero	1	1	2	2	2
3	audi	1	1	4	1	1
4	audi	1	1	4	1	3

	engine	location	wheelbase	carwidth	carheight	enginetype
cylindernumber \						
0	1	88.6	64.1	48.8	2	
4						
1	1	88.6	64.1	48.8	2	
4						
2	1	94.5	65.5	52.4	2	
6						
3	1	99.8	66.2	54.3	1	
4						
4	1	99.4	66.4	54.3	1	
1						

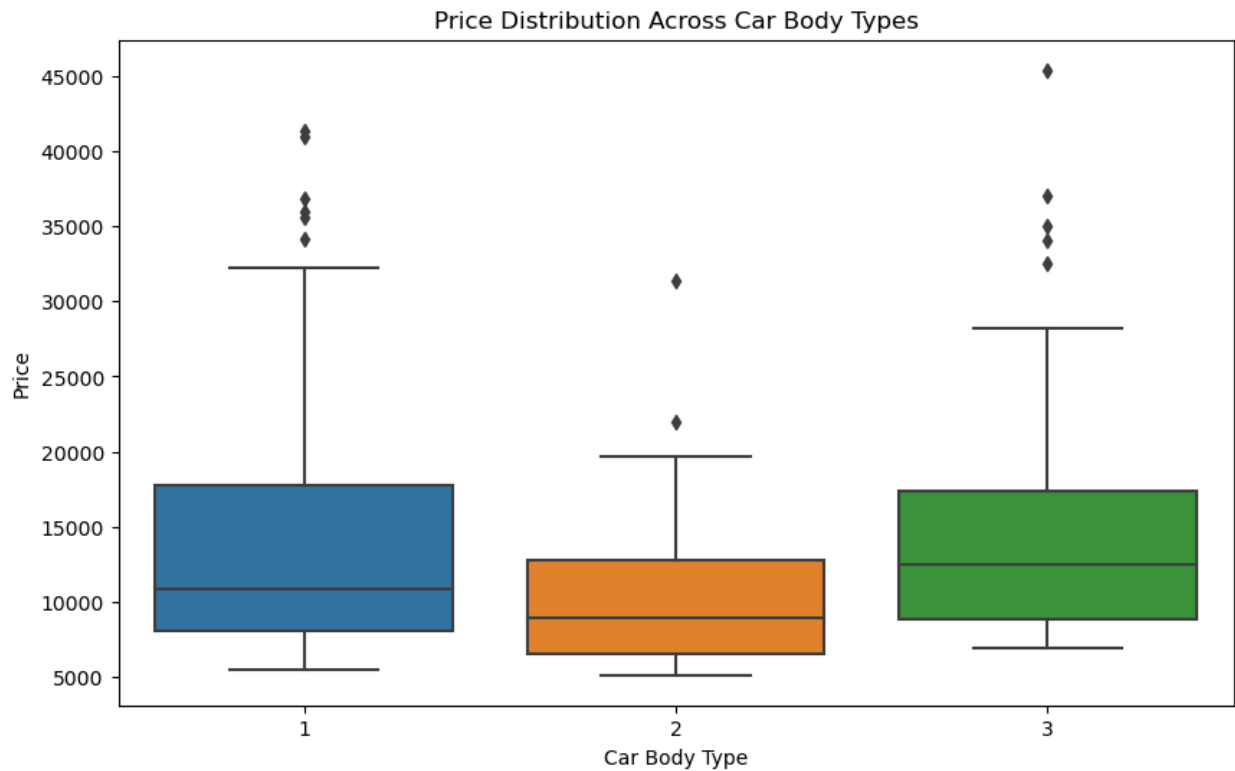
	enginesize	fuelsystem	boreratio	horsepower	price
0	130	1	3.47	111	13495.0
1	130	1	3.47	111	16500.0
2	152	1	2.68	154	16500.0
3	109	1	3.19	102	13950.0
4	136	1	3.19	115	17450.0

```
plt.figure(figsize=(10, 6))
sns.histplot(df['price'], kde=True, bins=20, color='green')
plt.title('Distribution of Car Prices')
plt.xlabel('Price')
plt.ylabel('Density')
plt.show()
```

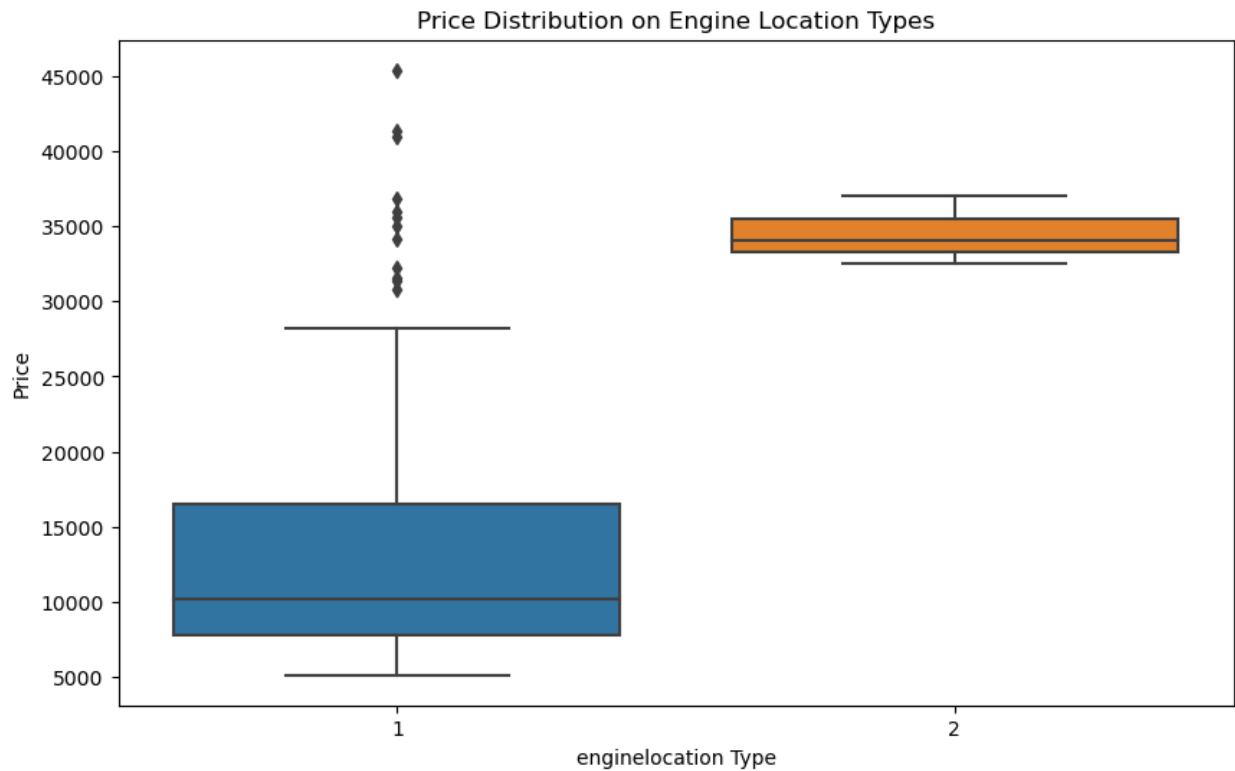




```
# Create a box plot
plt.figure(figsize=(10, 6))
sns.boxplot(data=df, x='carbody', y='price')
plt.xlabel('Car Body Type')
plt.ylabel('Price')
plt.title('Price Distribution Across Car Body Types')
plt.show()
```




```
# Create a box plot
plt.figure(figsize=(10, 6))
sns.boxplot(data=df, x='engine_location', y='price')
plt.xlabel('engine_location Type')
plt.ylabel('Price')
plt.title('Price Distribution on Engine Location Types')
plt.show()
```



```
# Select only numerical columns
numeric_df = df.select_dtypes(include=['float64', 'int64'])

plt.figure(figsize=(16, 8))
sns.heatmap(numeric_df.corr(), annot=True, cmap='coolwarm')
plt.title('Correlation Matrix Heatmap')
plt.show()
```



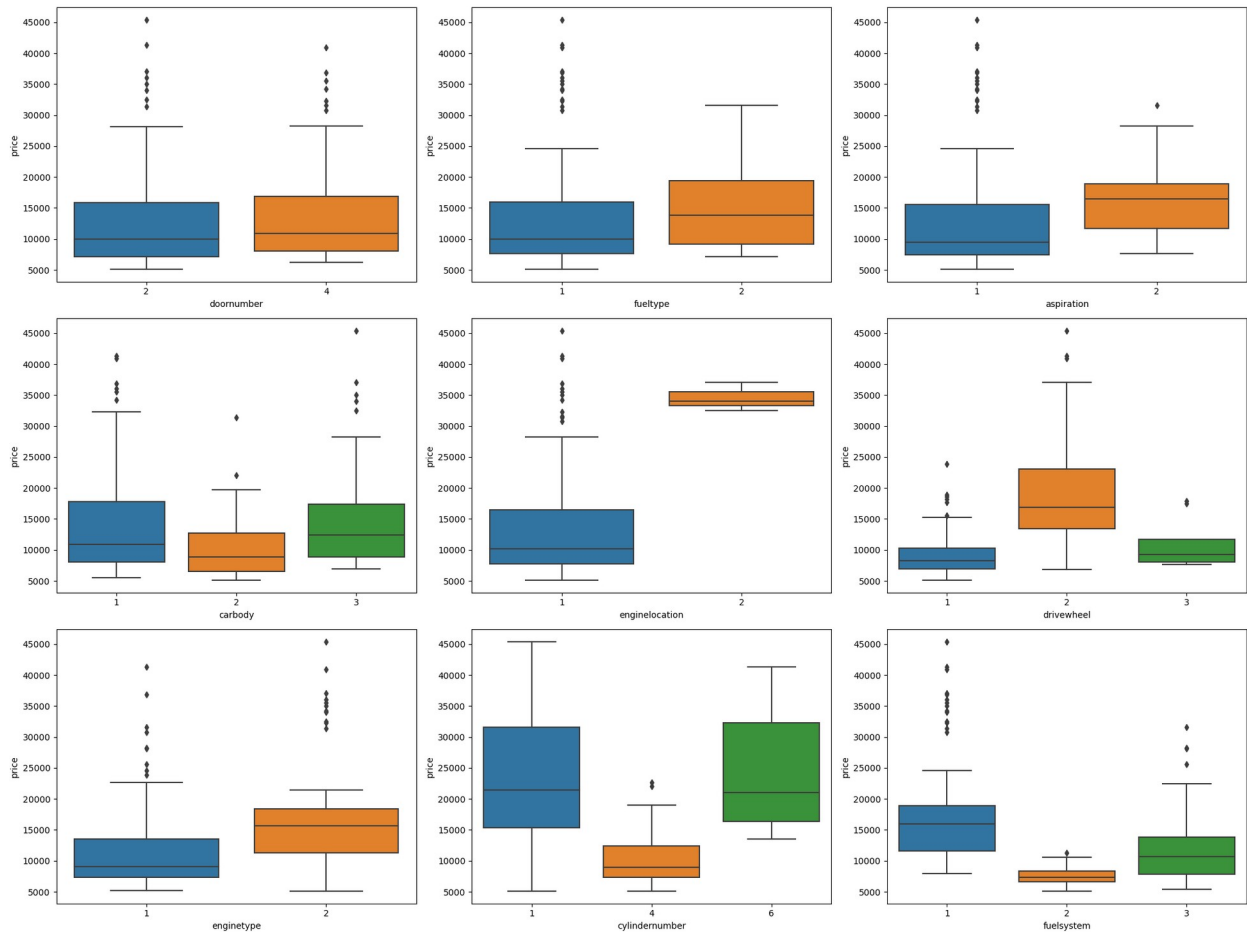
```
import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(20, 15))

# List of categorical columns
cat_columns = ['doornumber', 'fueltype', 'aspiration', 'carbody',
               'enginelocation',
               'drivewheel', 'enginetype', 'cylindernumber',
               'fuelsystem']

# Loop through each categorical column and create a boxplot
for i, col in enumerate(cat_columns, 1):
    plt.subplot(3, 3, i)
    sns.boxplot(x=col, y='price', data=df)

plt.tight_layout()
plt.show()
```



Diesel fueltype cars are comparatively expensive than the cars with fueltype as gas.

All the types of carbody is relatively cheaper as compared to convertible and hardtop carbody.

The cars with rear enginelocation are way expensive than cars with front enginelocation.

rwd drivewheel car price is always very high.

Enginetype ohcv comes into higher price range cars.

The price of car is directly proportional to no. of cylinders in most cases.

```
# Import necessary libraries
from sklearn.model_selection import train_test_split # Import
train_test_split function for splitting data
from sklearn import preprocessing # Import preprocessing module for
data preprocessing tasks
from sklearn.metrics import mean_absolute_error
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.preprocessing import StandardScaler
```

```

# Select numeric columns from the DataFrame
data_numeric = df.select_dtypes(include=['float64', 'int64'])

# Define independent variables (X) and dependent variable (y)
X = data_numeric.drop(columns=['price']) # Independent variables
excluding the dependent variables
y = data_numeric['price'] # Dependent variables

data_numeric = df.select_dtypes(include=['float64', 'int64'])

# Define independent variables (X) and dependent variable (y)
X = data_numeric.drop(columns=['price']) # Independent variables
y = data_numeric['price'] # Dependent variable

# Split the data into training and testing sets
X_train, X_test, y_train, y_test =
train_test_split(X,y,random_state=0)
model = LinearRegression()
model.fit(X_train,y_train)

y_pred = model.predict(X_test)
MSE = mean_squared_error(y_test,y_pred)
print(f"Mean Squared Error: {MSE}")
rmse = np.sqrt(MSE)
print(f"Root Mean Squared Error: {rmse}")
r2 = r2_score(y_test,y_pred) * 100
print(f"R2 Score: {r2}")

Mean Squared Error: 12240940.842944507
Root Mean Squared Error: 3498.7055953515874
R2 Score: 83.58419628988368

```

Model has an MSE of approximately 11,874,973.25, an RMSE(Root Mean Squared Error) of approximately 3,446.01, and an  $R^2$  score of approximately 84.07%. These values suggest that your model performs reasonably well, with relatively low error and a good fit to the data. However, further analysis and comparison with alternative models may be warranted to assess the model's performance comprehensively.

```

from sklearn.preprocessing import StandardScaler

# Define feature matrix A and target vector b
A = data_numeric.drop("price", axis=1)
b = data_numeric["price"]
# Scale the features using StandardScaler
scaler = StandardScaler()
A = scaler.fit_transform(A)

# Split the scaled data into training and testing sets
X_train_s, X_test_s, y_train_s, y_test_s = train_test_split(A, b,
test_size=0.30, random_state=40)

```

```

# Print the shapes of the scaled data
print("X_train_scaler:", X_train_s.shape)
print("X_test_scaler:", X_test_s.shape)
print("Y_train_scaler:", y_train_s.shape)
print("Y_test_scaler:", y_test_s.shape)

X_train_scaler: (143, 15)
X_test_scaler: (62, 15)
Y_train_scaler: (143,)
Y_test_scaler: (62,)

model.fit(X_train_s,y_train_s)
y_pred = model.predict(X_test_s)

mse = mean_squared_error(y_test_s,y_pred)
print(f"Mean Squared Error: {mse}")
rmse = np.sqrt(mse)
print(f"Root Mean Squared Error: {rmse}")
r2 = r2_score(y_test_s,y_pred) * 100
print(f"R2 Score: {r2}")

Mean Squared Error: 5784164.607479063
Root Mean Squared Error: 2405.029024248785
R2 Score: 81.75755006601055

```

Here, it shows that the given mean sum of square show that the difference between our model predictiona and actual value. And R-squared of value 81.75 shows our explanatory variables explain about 81% to predict the model shows good models measure fitness.